

Lecture 4-5A

Problem Solving by Searching

State space graphs and uninformed search

Textbook: Artificial Intelligence, A Modern Approach, IV edition

Authors: Russell and Norvig

Reading material: Chapter 3, sections 3.1 to 3.4, pp. 63-84

© Professor Arvind Bansal

1


Outline

- ▶ State space graph
 - ▶ **state** is a set of variable-value pairs; graph-nodes are states
 - ▶ A move changes the value of a variable and transitions to new state
 - ▶ edge is a change in one or more values associated with variables
 - ▶ **goal** is when state meets the **final condition**
 - ▶ Components of a state-space graph: **initial state, set of states, moves that change one or more states, goals**
 - ▶ There is a cost associated with making a move; an edge may have varying weight
 - ▶ Total cost is sum of the cost of the path from initial state to the goal state
- ▶ Search strategies
 - ▶ exhaustive/blind search: depth first; breadth-first and their variations
- ▶ **issues in graph-based search**
 - ▶ how to prune the search space without sacrificing a solution

© Professor Arvind Bansal

2

State Space Graph




© Professor Arvind Bansal

3

AI Search

- ▶ Plays an important role in AI
- ▶ Key idea includes
 - "mathematically sound estimate of distance of the current state to the goal state"
 - "prune search space and remain focused"
 - smaller the difference between estimated distance and minimal distance, better pruning is achieved
- ▶ Applicability and advantages
 - ▶ real-world problems that require **exponential time**
 - ▶ handling very large search space where algorithmic approaches will consume unrealistic time



© Professor Arvind Bansal

4

AI Search Problems

- ▶ Finding the best route from one node to another node towards a goal state based upon a well-defined notion such as
 - ▶ minimum time / minimum distance / minimum cost
- ▶ Finding the best strategy to win a game / solve a puzzle
- ▶ Finding the solution to a logical query
- ▶ Finding the solution under constrained conditions, such as
 - ▶ laying the printed circuit board and VLSI design without jumpers
 - ▶ putting the fixed number of workers on a complex task
 - ▶ carrying maximum passengers with minimum planes
 - ▶ adjusting traffic lights time to minimize delay and congestion
 - ▶ adjusting electric grid to minimize power outages

© Professor Arvind Bansal

5

Heuristic Search Strategy


- ▶ Identify the problem
- ▶ Abstract and formulate the problem
 - ▶ extract the features for identifying the variables to form states
 - ▶ identify set of states, initial state, final state (goal state), and sets of incremental moves to take to next state
- ▶ Start with the initial state
- ▶ At every node, test if the 'goal state' has been reached
- ▶ Find the sequence of moves that will
 - ▶ take from **initial state to goal state** in a focused manner using heuristics functions
 - ▶ test and remove cyclic traversal – archive visited nodes in a set and perform membership test to avoid visiting visited nodes
 - ▶ minimize the total cost according to some criteria
 - ▶ search in reasonable time and memory as desired



© Professor Arvind Bansal

6

State Space Graph




© Professor Arvind Bansal

7

Example: Sliding Puzzle

- ▶ States:
 - ▶ description of the eight tiles and location of the blank tile
- ▶ Successor function:
 - ▶ generates intermediate states using one of four moves: {*left*, *right*, *up*, *down*}
- ▶ Goal test:
 - ▶ checks if the current state matches the final state
- ▶ Path cost:
 - ▶ each move is unit cost
 - ▶ total cost = number of moves




© Professor Arvind Bansal

8

Example: Eight Queens

- Place eight queens on a chess board such that no queen can attack another queen
 - naïve number of states: $64 * 63 * 62 * 61 * 60 * 59 * 58 * 57 = 1.78 * 10^{14}$
- Intelligent formulation with pruning
 - place one queen at a time
 - do not place a queen in square if it is attacked by other queens
 - new number of states is around $8 * 5 * 4 * 3 * 2 * 2 * 1 * 2 = 3840 \text{ states}$



© Professor Arvind Bansal

9

Traversal Planning

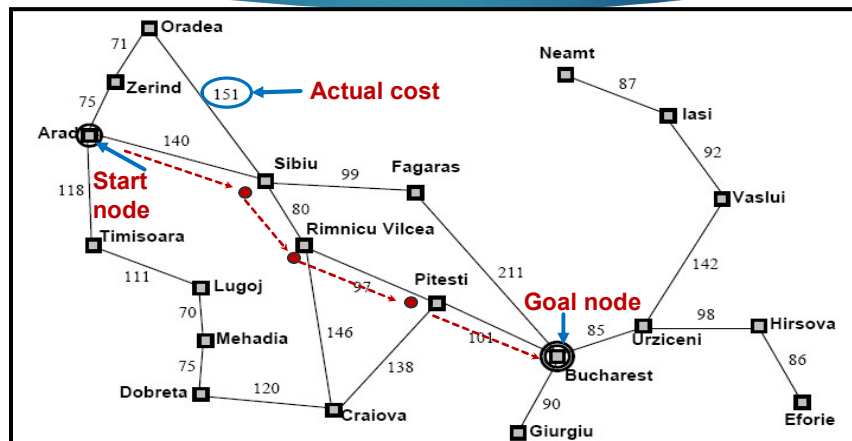


Figure taken from the book "Artificial Intelligence – A modern Approach by Russell and Norvig

10

Search Strategies


- ▶ Exhaustive search (blind search)
 - ▶ uninformed no cost estimated used to reach goal state
 - ▶ **Breadth-first, uniform-cost, depth-first, limited depth-first, iterative deepening**
- ▶ Search space pruning (informed search using heuristics)
 - ▶ focused movement towards the goal
 - ▶ uses a heuristic cost-function that estimates of distance from the current state to the goal state using mathematical abstraction
 - ▶ total cost = actual cost of traversal from initial state to next state + estimated cost from next state to goal state
 - ▶ **assumption: actual cost from current state to next state is known**
- ▶ Two major types of heuristic searches called **best first search**
 - Greedy best first**, estimated cost from the current state to the goal state is minimal
 - A* algorithm: total cost** from any traversed node to goal-state is minimal. Hence, **A* always finds the optimal path (at the cost of additional memory)**

© Professor Arvind Bansal

11

Terminology Used

- ▶ Complete : search finds out all the existing goal states / final states
- ▶ Incomplete: search misses out at least one goal state / final state
- ▶ Optimal search
 - ▶ finds out the least-cost solution by expanding the right nodes.
- ▶ Heuristics function: estimates distance from current state to goal state
- ▶ Pruning
 - ▶ search does not traverse a subset of nodes that either has a higher total cost or do not lead to goal-state based upon heuristic-function estimates
- ▶ Frontier (or fringe) – set of leaf nodes of the **currently traversed part** of the tree



© Professor Arvind Bansal

12

General Tree Search Strategy

```

function TREE-SEARCH(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion then return failure
    choose one node from frontier according to the current strategy
    if the node contains a goal state then return the corresponding solution
    else expand the node and add the resulting nodes to the search tree
  end

```

Note: Tree is a subset of directed acyclic graphs (DAGs). The algorithm can be generalized to any DAG traversal.

Algorithm taken from the book "Artificial Intelligence – A modern Approach by Russell and Norwig

13

Generic Tree-search Algorithm

Initially fringe is empty

```



function TREE-SEARCH(problem, fringe) returns a solution, or failure
  fringe ← INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node ← REMOVE-FRONT(fringe)
    If goal_test(state(node)) then return(node)
    fringe ← INSERTALL(EXPAND(node, problem), fringe)

```

```

function EXPAND(node, problem) returns a set of nodes
  successors ← the empty set
  for each action, result in SUCCESSOR-FN[problem](STATE[node]) do
    s ← a new NODE
    PARENT-NODE[s] ← node; ACTION[s] ← action; STATE[s] ← result
    PATH-COST[s] ← PATH-COST[node] + STEP-COST(node, action, s)
    DEPTH[s] ← DEPTH[node] + 1
    add s to successors
  return successors

```



Algorithm taken from the book "Artificial Intelligence – A modern Approach by Russell and Norwig

14

Search Strategies


- ▶ A search strategy is defined by picking the order of node expansion
- ▶ Strategies are evaluated using the following criteria:
 - ▶ **completeness**: does it always find a solution if one exists?
 - ▶ **time complexity**: number of nodes generated and traversed
 - ▶ **space complexity**: maximum number of nodes in the fringe queue
 - ▶ **optimality**: does it always find a least-cost solution?
- ▶ Time and space complexity are measured in terms of
 - ▶ **b**: maximum branching factor of the search tree
 - ▶ **d**: depth of the least-cost solution
 - ▶ **m**: maximum depth of the state-space (may be **very large**)

© Professor Arvind Bansal

15

Exhaustive / Blind Search

- ▶ Does not use heuristic function. Traverses potentially every node.
- ▶ Breadth-first search (BFS)
 - ▶ search all nodes level by level ; uses a fringe queue to store all the children of the expanded unexplored node
- ▶ Uniform cost search
 - ▶ a variation of breadth-first search that picks the shortest-path node during node expansion – **replace queue by a priority queue**
- ▶ Depth-first search (DFS)
 - ▶ search the leftmost children first (**similar to pre-order traversal**)
 - ▶ uses stack to store deferred right siblings
- ▶ Depth-limited search (DLS)
 - ▶ a variation of depth-first search that limits the depth being searched
- ▶ Iterative deepening search (IDS or IDDFS)
 - ▶ iterative deepening of depth and performing depth-first search




© Professor Arvind Bansal

16

Breadth-first Search

Assumption: Goal state is the leaf node of the tree

1. Put the root node in the queue;
2. **Loop** % traverse level wise
 1. pick the next node from the queue;
 2. **If** it is a leaf node **then**
 If node is a goal-state **then**
 return(node) and **exit**;
 3. **else if** it is a non-leaf node **then**
 find all the children of the node,
 and insert them in the queue; }
3. **until** the queue is empty
5. If queue is empty then return **failure**;



© Professor Arvind Bansal

17

Properties of Breadth-first Search

- ▶ Complete? Yes (if branching factor b is finite)
- ▶ Time? $1 + b + b^2 + b^3 + \dots + b^d = O(b^{d+1})$ where d is depth of the tree
- ▶ Space? $O(b^{d+1})$ (keeps every unexpanded node in the queue)
- ▶ Optimal? Yes (if cost = 1 per step)

© Professor Arvind Bansal

18

Uniform-cost Search

- ▶ Expand **least-cost unexpanded node**
 - ▶ use priority queue with **least path cost from root to the current node**
- ▶ Implementation:
 - ▶ **fringe = queue sorted in ascending order**
 - ▶ **goal-node test done before expansion and not after generation**
- ▶ Becomes breadth-first if step costs all equal
- ▶ Complete? Yes
- ▶ Time? # $O(b^{\lceil C^*/\epsilon \rceil})$; C^* is cost of the optimal solution
- ▶ Space? # $O(b^{\lceil C^*/\epsilon \rceil})$
- ▶ Optimal? Yes – nodes expanded in increasing order of cost




Figure taken from the book "Artificial Intelligence – A modern Approach by Russell and Norvig, 4th edition

19

Uniform Cost Search Algorithm

```


function UNIFORM-COST-SEARCH(problem) returns a solution, or failure
  node ← a node with STATE = problem.INITIAL-STATE, PATH-COST = 0
  frontier ← a priority queue ordered by PATH-COST, with node as the only element
  explored ← an empty set
  loop do
    if EMPTY?(frontier) then return failure
    node ← POP(frontier) /* chooses the lowest-cost node in frontier */
    if problem.GOAL-TEST(node.STATE) then return SOLUTION(node)
    add node.STATE to explored
    for each action in problem.ACTIONS(node.STATE) do
      child ← CHILD-NODE(problem, node, action)
      if child.STATE is not in explored or frontier then
        frontier ← INSERT(child, frontier)
      else if child.STATE is in frontier with higher PATH-COST then
        replace that frontier node with child
  
```

Algorithm taken from the book "Artificial Intelligence – A modern Approach by Russell and Norvig

20

Depth-first Search

- ▶ $cn = \text{root-node}$; $\text{goal-found} = \text{false}$;
- ▶ **loop**
 - ▶ **if** (cn is a leaf node **&&** $\text{goal-state}(cn)$)
 $\text{goal-found} = \text{true}$; **return**(cn) and **exit**
 - ▶ **elseif** ($\text{leaf}(cn)$ **&&** $\neg \text{empty}(\text{stack})$)
 $cn = \text{pop}(\text{stack})$;
 - ▶ **elseif** $\neg \text{has-left-child}(cn)$
 $cn = \text{right-child}(cn)$;
 - ▶ **else** $\text{push}(\text{stack}, \text{right-sibling}(cn))$;
 - ▶ $cn = \text{left-child}(cn)$;
- ▶ **until** ($(\text{empty}(\text{stack}) \text{ \&\& } \text{leaf}(cn))$)
- ▶ **if** ($\text{empty}(\text{stack})$) **&&** $\text{leaf}(cn)$) **return failure**.



© Professor Arvind Bansal

21

Properties of Depth-first Search

- ▶ Complete? No: fails in infinite-depth spaces, spaces with loops
 - ▶ avoid repeated states along path using a lookup table of visited nodes
 - ▶ Generally, stops after finding the first solution
- ▶ Worst case Time? $O(b^m)$. Excessive if m is much larger than d
 - ▶ generally faster than breadth-first because all branches are not explored
 - ▶ Best case time: $O(bm)$ – when the goal is near the leftmost leaf node
- ▶ Space? $O(bm)$, i.e., linear stack space guided by depth of the tree
- ▶ Optimal? No, solution may be at higher level on the right side

© Professor Arvind Bansal

22

Depth-limited Search

- ▶ Depth-first search with depth-limit $L < \text{total depth of the tree}$,
 - ▶ nodes beyond depth L are not traversed
- ▶ Alleviates the problem of non-termination with indefinite branches
- ▶ Improves the search time
- ▶ **Problem with completeness**
 - ▶ can not find a goal-state if the goal-state is beyond depth L

© Professor Arvind Bansal

23

Depth Limited Search Recursive Algorithm

```

function DEPTH-LIMITED-SEARCH(problem, limit) returns a solution, or failure/cutoff
  return RECURSIVE-DLS(MAKE-NODE(problem.INITIAL-STATE), problem, limit)

function RECURSIVE-DLS(node, problem, limit) returns a solution, or failure/cutoff
  if problem.GOAL-TEST(node.STATE) then return SOLUTION(node)
  else if limit = 0 then return cutoff
  else
    cutoff_occurred?  $\leftarrow$  false
    for each action in problem.ACTIONS(node.STATE) do
      child  $\leftarrow$  CHILD-NODE(problem, node, action)
      result  $\leftarrow$  RECURSIVE-DLS(child, problem, limit - 1)
      if result = cutoff then cutoff_occurred?  $\leftarrow$  true
      else if result  $\neq$  failure then return result
    if cutoff_occurred? then return cutoff else return failure
  
```

Algorithm taken from the book "Artificial Intelligence – A modern Approach by Russell and Norvig


24

Iterative Deepening Search (IDS)

- ▶ Iteratively perform depth-limited search
 - ▶ Level = 0, level = 1, level = 2 ... duplicates traversal of previous levels
- ▶ Advantages
 - ▶ Uses stack reducing space requirement to $O(bd)$
 - ▶ Traversed nodes limited by $O(b^d)$
 - ▶ much less than breadth-first search
 - ▶ Optimal like breadth-first search



© Professor Arvind Bansal



25

IDS Algorithm

function ITERATIVE-DEEPENING-SEARCH(*problem*) **returns** a solution node or *failure*
for *depth* = 0 **to** ∞ **do**
 result \leftarrow DEPTH-LIMITED-SEARCH(*problem*, *depth*)
 if *result* \neq *cutoff* **then return** *result*





Figure taken from the book "Artificial Intelligence – A modern Approach by Russell and Norvig

26

Bidirectional Search

- ▶ Start in both the directions
 - ▶ **goal-state \rightarrow initial state** and
 - initial-state \rightarrow goal-state**
 - Meet somewhere halfway**
- ▶ Expand until the two fringes intersect
- ▶ Complexity: $O(b^{d/2})$
- ▶ Disadvantages
 - ▶ lot of memory needed to keep fringes in memory for intersection
 - ▶ overhead of finding intersection after every move
 - ▶ backward moves from goal-states not always easy to figure out



© Professor Arvind Bansal

27

Comparing Uninformed Search Techniques

Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening	Bidirectional (if applicable)
Complete?	Yes ¹	Yes ^{1,2}	No	No	Yes ¹	Yes ^{1,4}
Optimal cost?	Yes ³	Yes	No	No	Yes ³	Yes ^{3,4}
Time	$O(b^d)$	$O(b^{1+\lceil C^*/\epsilon \rceil})$	$O(b^m)$	$O(b^\ell)$	$O(b^d)$	$O(b^{d/2})$
Space	$O(b^d)$	$O(b^{1+\lceil C^*/\epsilon \rceil})$	$O(bm)$	$O(b\ell)$	$O(bd)$	$O(b^{d/2})$

Table taken from the book "Artificial Intelligence – A modern Approach by Russell and Norvig

28

Graph-based Search

```

function GRAPH-SEARCH(problem, fringe) returns a solution, or failure
  closed ← an empty set ; fringe ← an empty set;
  fringe ← INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node ← REMOVE-FRONT(fringe)
    if GOAL-TEST[problem](STATE[node]) then return SOLUTION(node)
    if STATE[node] is not in closed then
      add STATE[node] to closed
      fringe ← INSERTALL(EXPAND(node, problem), fringe)


```

Algorithm taken from the book "Artificial Intelligence – A modern Approach by Russell and Norwig

29

Handling Repeated States

- ▶ Caused by
 - ▶ cycles in the search path
 - ▶ multiple moves between two states giving rise to duplications
- ▶ Causes looping and/or exponential explosion of states
- ▶ Solution
 - ▶ keep history of traversed nodes
 - ▶ **perform membership test** using hashing for fast search



© Professor Arvind Bansal

30